

Earth Quake Damage Prediction Model Using Machine Learning

Importing essential libraries

```
In [1]: ▶ import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as ss
from scipy.stats import chi2_contingency
from scipy.stats import chi2
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import homogeneity_score
from sklearn.metrics import silhouette_score
from sklearn.metrics import classification_report, f1_score, precision_score, recall_sc

#To Ignore Warnings
import warnings
warnings.filterwarnings('ignore')
```

Importing and understanding our dataset

```
In [2]: ▶ train=pd.read_csv('D:\Major Project\Dataset/train.csv')
```

Shape of dataset

```
In [3]: ▶ train.shape
```

```
Out[3]: (631761, 14)
```

Printing out a few columns

```
In [4]: ▶ train.head()
```

```
Out[4]:
```

	area_assesed	building_id	damage_grade	district_id	has_geotechnical_risk	has_geotechnical_risk_f
0	Both	24385bfd2a2	Grade 4	24		0.0
1	Both	405d1bbebbf	Grade 2	44		0.0
2	Both	351d9bc71f6	Grade 1	36		0.0
3	Building removed	2be3a971166	Grade 5	30		0.0
4	Both	34c7d073ea6	Grade 3	36		0.0

```
In [5]: train.columns
```

```
Out[5]: Index(['area_assessed', 'building_id', 'damage_grade', 'district_id',  
              'has_geotechnical_risk', 'has_geotechnical_risk_fault_crack',  
              'has_geotechnical_risk_flood', 'has_geotechnical_risk_land_settlement',  
              'has_geotechnical_risk_landslide', 'has_geotechnical_risk_liquefaction',  
              'has_geotechnical_risk_other', 'has_geotechnical_risk_rock_fall',  
              'has_repair_started', 'vdcmun_id'],  
              dtype='object')
```

```
In [6]: len(train.columns)
```

```
Out[6]: 14
```

```
In [7]: train.isnull().sum()
```

```
Out[7]: area_assessed      0  
        building_id      0  
        damage_grade      0  
        district_id      0  
        has_geotechnical_risk      0  
        has_geotechnical_risk_fault_crack      0  
        has_geotechnical_risk_flood      0  
        has_geotechnical_risk_land_settlement      0  
        has_geotechnical_risk_landslide      0  
        has_geotechnical_risk_liquefaction      0  
        has_geotechnical_risk_other      0  
        has_geotechnical_risk_rock_fall      0  
        has_repair_started    33417  
        vdcmun_id      0  
        dtype: int64
```

Using of other datasets

```
In [8]: owner = pd.read_csv('D:\Major Project\Dataset\Building_Ownership_Use.csv')
structure = pd.read_csv('D:\Major Project\Dataset\Building_Structure.csv')
owner.info()
structure.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1052948 entries, 0 to 1052947
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	building_id	1052948 non-null	object
1	district_id	1052948 non-null	int64
2	vdcmun_id	1052948 non-null	int64
3	ward_id	1052948 non-null	int64
4	legal_ownership_status	1052948 non-null	object
5	count_families	1052946 non-null	float64
6	has_secondary_use	1052938 non-null	float64
7	has_secondary_use_agriculture	1052948 non-null	int64
8	has_secondary_use_hotel	1052948 non-null	int64
9	has_secondary_use_rental	1052948 non-null	int64
10	has_secondary_use_institution	1052948 non-null	int64
11	has_secondary_use_school	1052948 non-null	int64
12	has_secondary_use_industry	1052948 non-null	int64
13	has_secondary_use_health_post	1052948 non-null	int64
14	has_secondary_use_gov_office	1052948 non-null	int64
15	has_secondary_use_use_police	1052948 non-null	int64
16	has_secondary_use_other	1052948 non-null	int64

```
dtypes: float64(2), int64(13), object(2)
```

```
memory usage: 136.6+ MB
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1052948 entries, 0 to 1052947
Data columns (total 29 columns):
```

#	Column	Non-Null Count	Dtype
0	building_id	1052948 non-null	object
1	district_id	1052948 non-null	int64
2	vdcmun_id	1052948 non-null	int64
3	ward_id	1052948 non-null	int64
4	count_floors_pre_eq	1052948 non-null	int64
5	count_floors_post_eq	1052948 non-null	int64
6	age_building	1052948 non-null	int64
7	plinth_area_sq_ft	1052948 non-null	int64
8	height_ft_pre_eq	1052948 non-null	int64
9	height_ft_post_eq	1052948 non-null	int64
10	land_surface_condition	1052948 non-null	object
11	foundation_type	1052948 non-null	object
12	roof_type	1052948 non-null	object
13	ground_floor_type	1052948 non-null	object
14	other_floor_type	1052948 non-null	object
15	position	1052947 non-null	object
16	plan_configuration	1052947 non-null	object
17	has_superstructure_adobe_mud	1052948 non-null	int64
18	has_superstructure_mud_mortar_stone	1052948 non-null	int64
19	has_superstructure_stone_flag	1052948 non-null	int64
20	has_superstructure_cement_mortar_stone	1052948 non-null	int64
21	has_superstructure_mud_mortar_brick	1052948 non-null	int64
22	has_superstructure_cement_mortar_brick	1052948 non-null	int64
23	has_superstructure_timber	1052948 non-null	int64
24	has_superstructure_bamboo	1052948 non-null	int64
25	has_superstructure_rc_non_engineered	1052948 non-null	int64
26	has_superstructure_rc_engineered	1052948 non-null	int64
27	has_superstructure_other	1052948 non-null	int64
28	condition_post_eq	1052948 non-null	object

```
dtypes: int64(20), object(9)
```

```
memory usage: 233.0+ MB
```

Combining of other datasets for preprocessing

```
In [9]: ► combine = pd.merge(owner,structure, on='building_id')
trainfinal = pd.merge(combine,train, on = 'building_id')
```

Train Data before preprocessing

```
In [10]: ► trainfinal.columns,len(trainfinal.columns)
```

```
Out[10]: (Index(['building_id', 'district_id_x', 'vdcmun_id_x', 'ward_id_x',
'legal_ownership_status', 'count_families', 'has_secondary_use',
'has_secondary_use_agriculture', 'has_secondary_use_hotel',
'has_secondary_use_rental', 'has_secondary_use_institution',
'has_secondary_use_school', 'has_secondary_use_industry',
'has_secondary_use_health_post', 'has_secondary_use_gov_office',
'has_secondary_use_use_police', 'has_secondary_use_other',
'district_id_y', 'vdcmun_id_y', 'ward_id_y', 'count_floors_pre_eq',
'count_floors_post_eq', 'age_building', 'plinth_area_sq_ft',
'height_ft_pre_eq', 'height_ft_post_eq', 'land_surface_condition',
'foundation_type', 'roof_type', 'ground_floor_type', 'other_floor_type',
'position', 'plan_configuration', 'has_superstructure_adobe_mud',
'has_superstructure_mud_mortar_stone', 'has_superstructure_stone_flag',
'has_superstructure_cement_mortar_stone',
'has_superstructure_mud_mortar_brick',
'has_superstructure_cement_mortar_brick', 'has_superstructure_timber',
'has_superstructure_bamboo', 'has_superstructure_rc_non_engineered',
'has_superstructure_rc_engineered', 'has_superstructure_other',
'condition_post_eq', 'area_assesed', 'damage_grade', 'district_id',
'has_geotechnical_risk', 'has_geotechnical_risk_fault_crack',
'has_geotechnical_risk_flood', 'has_geotechnical_risk_land_settlement',
'has_geotechnical_risk_landslide', 'has_geotechnical_risk_liquefaction',
'has_geotechnical_risk_other', 'has_geotechnical_risk_rock_fall',
'has_repair_started', 'vdcmun_id'],
dtype='object'),
58)
```

```
In [11]: ▶ trainfinal.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 631761 entries, 0 to 631760
Data columns (total 58 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   building_id                                                            631761 non-null object
1   district_id_x                                                          631761 non-null int64
2   vdcmun_id_x                                                            631761 non-null int64
3   ward_id_x                                                              631761 non-null int64
4   legal_ownership_status                                                631761 non-null object
5   count_families                                                         631760 non-null float64
6   has_secondary_use                                                      631761 non-null float64
7   has_secondary_use_agriculture                                          631761 non-null int64
8   has_secondary_use_hotel                                                631761 non-null int64
9   has_secondary_use_rental                                              631761 non-null int64
10  has_secondary_use_institution                                          631761 non-null int64
11  has_secondary_use_school                                              631761 non-null int64
12  has_secondary_use_industry                                             631761 non-null int64
13  has_secondary_use_health_post                                         631761 non-null int64
14  has_secondary_use_gov_office                                          631761 non-null int64
15  has_secondary_use_use_police                                          631761 non-null int64
16  has_secondary_use_other                                               631761 non-null int64
17  district_id_y                                                          631761 non-null int64
18  vdcmun_id_y                                                            631761 non-null int64
19  ward_id_y                                                              631761 non-null int64
20  count_floors_pre_eq                                                    631761 non-null int64
21  count_floors_post_eq                                                  631761 non-null int64
22  age_building                                                           631761 non-null int64
23  plinth_area_sq_ft                                                      631761 non-null int64
24  height_ft_pre_eq                                                       631761 non-null int64
25  height_ft_post_eq                                                      631761 non-null int64
26  land_surface_condition                                                 631761 non-null object
27  foundation_type                                                        631761 non-null object
28  roof_type                                                              631761 non-null object
29  ground_floor_type                                                      631761 non-null object
30  other_floor_type                                                       631761 non-null object
31  position                                                                631761 non-null object
32  plan_configuration                                                     631761 non-null object
33  has_superstructure_adobe_mud                                           631761 non-null int64
34  has_superstructure_mud_mortar_stone 631761 non-null int64
35  has_superstructure_stone_flag                                          631761 non-null int64
36  has_superstructure_cement_mortar_stone 631761 non-null int64
37  has_superstructure_mud_mortar_brick 631761 non-null int64
38  has_superstructure_cement_mortar_brick 631761 non-null int64
39  has_superstructure_timber                                              631761 non-null int64
40  has_superstructure_bamboo                                              631761 non-null int64
41  has_superstructure_rc_non_engineered 631761 non-null int64
42  has_superstructure_rc_engineered   631761 non-null int64
43  has_superstructure_other                                               631761 non-null int64
44  condition_post_eq                                                      631761 non-null object
45  area_assesed                                                           631761 non-null object
46  damage_grade                                                           631761 non-null object
47  district_id                                                            631761 non-null int64
48  has_geotechnical_risk                                                  631761 non-null float64
49  has_geotechnical_risk_fault_crack 631761 non-null int64
50  has_geotechnical_risk_flood                                            631761 non-null int64
51  has_geotechnical_risk_land_settlement 631761 non-null int64
52  has_geotechnical_risk_landslide   631761 non-null int64
53  has_geotechnical_risk_liquefaction 631761 non-null int64
54  has_geotechnical_risk_other                                             631761 non-null int64
55  has_geotechnical_risk_rock_fall   631761 non-null int64
56  has_repair_started                                                     598344 non-null float64
57  vdcmun_id                                                              631761 non-null int64
dtypes: float64(4), int64(42), object(12)
memory usage: 284.4+ MB
```

In [12]: `trainfinal.isnull().sum()`

```
Out[12]: building_id                0
district_id_x                    0
vdcmun_id_x                     0
ward_id_x                       0
legal_ownership_status          0
count_families                  1
has_secondary_use               0
has_secondary_use_agriculture   0
has_secondary_use_hotel         0
has_secondary_use_rental        0
has_secondary_use_institution   0
has_secondary_use_school        0
has_secondary_use_industry      0
has_secondary_use_health_post   0
has_secondary_use_gov_office    0
has_secondary_use_use_police    0
has_secondary_use_other         0
district_id_y                   0
vdcmun_id_y                     0
ward_id_y                       0
count_floors_pre_eq             0
count_floors_post_eq            0
age_building                    0
plinth_area_sq_ft              0
height_ft_pre_eq                0
height_ft_post_eq               0
land_surface_condition          0
foundation_type                 0
roof_type                       0
ground_floor_type               0
other_floor_type                0
position                        0
plan_configuration              0
has_superstructure_adobe_mud    0
has_superstructure_mud_mortar_stone 0
has_superstructure_stone_flag   0
has_superstructure_cement_mortar_stone 0
has_superstructure_mud_mortar_brick 0
has_superstructure_cement_mortar_brick 0
has_superstructure_timber       0
has_superstructure_bamboo       0
has_superstructure_rc_non_engineered 0
has_superstructure_rc_engineered 0
has_superstructure_other        0
condition_post_eq              0
area_assesed                   0
damage_grade                    0
district_id                     0
has_geotechnical_risk           0
has_geotechnical_risk_fault_crack 0
has_geotechnical_risk_flood     0
has_geotechnical_risk_land_settlement 0
has_geotechnical_risk_landslide 0
has_geotechnical_risk_liquefaction 0
has_geotechnical_risk_other     0
has_geotechnical_risk_rock_fall 0
has_repair_started              33417
vdcmun_id                       0
dtype: int64
```

```
In [13]: trainfinal.drop('has_repair_started',axis=1,inplace=True)
```

```
In [14]: trainfinal.drop(['vdcmun_id_y','district_id_y','ward_id_y','vdcmun_id','district_id
```

```
In [15]: features = list(trainfinal.columns)
```

Let's understand our columns better:

In [16]: ▶ trainfinal.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 631761 entries, 0 to 631760
Data columns (total 52 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   building_id                                                            631761 non-null  object
1   district_id_x                                                          631761 non-null  int64
2   vdcmun_id_x                                                            631761 non-null  int64
3   ward_id_x                                                              631761 non-null  int64
4   legal_ownership_status                                                631761 non-null  object
5   count_families                                                         631760 non-null  float64
6   has_secondary_use                                                      631761 non-null  float64
7   has_secondary_use_agriculture                                          631761 non-null  int64
8   has_secondary_use_hotel                                                631761 non-null  int64
9   has_secondary_use_rental                                              631761 non-null  int64
10  has_secondary_use_institution                                          631761 non-null  int64
11  has_secondary_use_school                                               631761 non-null  int64
12  has_secondary_use_industry                                              631761 non-null  int64
13  has_secondary_use_health_post                                          631761 non-null  int64
14  has_secondary_use_gov_office                                           631761 non-null  int64
15  has_secondary_use_use_police                                           631761 non-null  int64
16  has_secondary_use_other                                                631761 non-null  int64
17  count_floors_pre_eq                                                    631761 non-null  int64
18  count_floors_post_eq                                                   631761 non-null  int64
19  age_building                                                            631761 non-null  int64
20  plinth_area_sq_ft                                                      631761 non-null  int64
21  height_ft_pre_eq                                                       631761 non-null  int64
22  height_ft_post_eq                                                      631761 non-null  int64
23  land_surface_condition                                                  631761 non-null  object
24  foundation_type                                                         631761 non-null  object
25  roof_type                                                               631761 non-null  object
26  ground_floor_type                                                       631761 non-null  object
27  other_floor_type                                                        631761 non-null  object
28  position                                                                631761 non-null  object
29  plan_configuration                                                      631761 non-null  object
30  has_superstructure_adobe_mud                                            631761 non-null  int64
31  has_superstructure_mud_mortar_stone  631761 non-null  int64
32  has_superstructure_stone_flag                                           631761 non-null  int64
33  has_superstructure_cement_mortar_stone 631761 non-null  int64
34  has_superstructure_mud_mortar_brick  631761 non-null  int64
35  has_superstructure_cement_mortar_brick 631761 non-null  int64
36  has_superstructure_timber                                                631761 non-null  int64
37  has_superstructure_bamboo                                                631761 non-null  int64
38  has_superstructure_rc_non_engineered  631761 non-null  int64
39  has_superstructure_rc_engineered    631761 non-null  int64
40  has_superstructure_other                                                  631761 non-null  int64
41  condition_post_eq                                                       631761 non-null  object
42  area_assesed                                                            631761 non-null  object
43  damage_grade                                                            631761 non-null  object
44  has_geotechnical_risk                                                    631761 non-null  float64
45  has_geotechnical_risk_fault_crack   631761 non-null  int64
46  has_geotechnical_risk_flood        631761 non-null  int64
47  has_geotechnical_risk_land_settlement 631761 non-null  int64
48  has_geotechnical_risk_landslide     631761 non-null  int64
49  has_geotechnical_risk_liquefaction  631761 non-null  int64
50  has_geotechnical_risk_other         631761 non-null  int64
51  has_geotechnical_risk_rock_fall     631761 non-null  int64
dtypes: float64(3), int64(37), object(12)
memory usage: 255.5+ MB
```

Converting Nominal to Numeric for the purpose of classification(Target Variable)


```
In [17]: ► conversion = {'damage_grade' : {"Grade 1" : 1, "Grade 2" : 2, "Grade 3" : 3, "Grade 4" : 4}}
train_temp = pd.DataFrame()
train_temp['damage_grade'] = trainfinal['damage_grade']

train_temp.replace(conversion, inplace = True)
trainfinal['damage_grade'] = train_temp['damage_grade']
```

```
In [18]: ► cat = [c for c in trainfinal if trainfinal[c].dtypes == "object"]
cat.remove('building_id')
print(cat)

['legal_ownership_status', 'land_surface_condition', 'foundation_type', 'roof_type', 'ground_floor_type', 'other_floor_type', 'position', 'plan_configuration', 'condition_post_eq', 'area_assessed']
```

Chi-Square test on Each Types for Understanding dependency of Target with each type

```
In [19]: ► def ChiSquareTest(cat, res_train):

    for c in cat:
        print(c)
        tab = pd.crosstab(res_train['damage_grade'], res_train[c])
        stat, p, dof, expected = chi2_contingency(tab)
        print('dof=%d' % dof)
        prob = 0.95
        critical = chi2.ppf(prob, dof)
        print('probability=%.3f, critical=%.3f, stat=%.3f' % (prob, critical, stat))
        if abs(stat) >= critical:
            print('Dependent (reject H0)')
        else:
            print('Independent (fail to reject H0)')
        # interpret p-value
        alpha = 1.0 - prob
        print('significance=%.3f, p=%.3f' % (alpha, p))
        if p <= alpha:
            print('Dependent (reject H0)')
        else:
            print('Independent (fail to reject H0)')

    print(" ")
```

Performing Chi-Square test on Object Types for Understanding dependency of Target with each object type

In [20]: ► ChiSquareTest(cat,trainfinal)

```
legal_ownership_status
dof=12
probability=0.950, critical=21.026, stat=8113.932
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

land_surface_condition
dof=8
probability=0.950, critical=15.507, stat=1408.392
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

foundation_type
dof=16
probability=0.950, critical=26.296, stat=138103.743
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

roof_type
dof=8
probability=0.950, critical=15.507, stat=85099.545
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

ground_floor_type
dof=16
probability=0.950, critical=26.296, stat=107177.045
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

other_floor_type
dof=12
probability=0.950, critical=21.026, stat=93978.074
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

position
dof=12
probability=0.950, critical=21.026, stat=5026.896
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

plan_configuration
dof=36
probability=0.950, critical=50.998, stat=4672.139
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

condition_post_eq
dof=28
probability=0.950, critical=41.337, stat=1112041.847
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

area_assesed
dof=16
```

```
probability=0.950, critical=26.296, stat=429576.136
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)
```

```
In [21]: ► cat_binary = [c for c in trainfinal if len(trainfinal[c].unique()) == 2]
cat_binary
```

```
Out[21]: ['has_secondary_use',
'has_secondary_use_agriculture',
'has_secondary_use_hotel',
'has_secondary_use_rental',
'has_secondary_use_institution',
'has_secondary_use_school',
'has_secondary_use_industry',
'has_secondary_use_health_post',
'has_secondary_use_gov_office',
'has_secondary_use_use_police',
'has_secondary_use_other',
'has_superstructure_adobe_mud',
'has_superstructure_mud_mortar_stone',
'has_superstructure_stone_flag',
'has_superstructure_cement_mortar_stone',
'has_superstructure_mud_mortar_brick',
'has_superstructure_cement_mortar_brick',
'has_superstructure_timber',
'has_superstructure_bamboo',
'has_superstructure_rc_non_engineered',
'has_superstructure_rc_engineered',
'has_superstructure_other',
'has_geotechnical_risk',
'has_geotechnical_risk_fault_crack',
'has_geotechnical_risk_flood',
'has_geotechnical_risk_land_settlement',
'has_geotechnical_risk_landslide',
'has_geotechnical_risk_liquefaction',
'has_geotechnical_risk_other',
'has_geotechnical_risk_rock_fall']
```

Performing Chi-Square test on Binary Types for Understanding dependency of Target with each object type

```
In [22]: ► ChiSquareTest(cat_binary,trainfinal)

probability=0.950, critical=9.488, stat=3825.490
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

has_superstructure_rc_non_engineered
dof=4
probability=0.950, critical=9.488, stat=28652.127
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)

has_superstructure_rc_engineered
dof=4
probability=0.950, critical=9.488, stat=37327.720
Dependent (reject H0)
significance=0.050, p=0.000
Dependent (reject H0)
```

```
has_superstructure_other
```

Removal columns on which Target is not dependent(By Chi-Square Test)

```
In [23]: trainfinal.drop(['has_secondary_use_use_police', 'building_id'], axis = 1, inplace=True)
```

Converting Nominal to Numeric in preprocessed data for the purpose of classification

```
In [24]: cont = [c for c in trainfinal if len(trainfinal[c].unique()) > 15]
indices = 0,1,2,3
cont = [i for j, i in enumerate(cont) if j not in indices]
```

```
In [25]: train_copy = trainfinal
```

```
In [26]: train_copy['IsPrivate'] = (train_copy["legal_ownership_status"] == "Private") * 1
```

```
In [27]: train_copy['IsFlat'] = (train_copy["land_surface_condition"] == "Flat") * 1
train_copy['IsMudFoundation'] = (train_copy["foundation_type"] == "Mud mortar-Stone/
train_copy['IsBambooRoofLight'] = (train_copy["roof_type"] == "Bamboo/Timber-Light r
train_copy['IsFloorTypeMud'] = (train_copy["ground_floor_type"] == "Mud") * 1
train_copy['OtherFloorTypeMud'] = (train_copy["other_floor_type"] == "Timber/Bamboo-
train_copy['IsNotAttached'] = (train_copy["position"] == "Not attached") * 1
train_copy['IsPlanConfigRectangular'] = (train_copy["plan_configuration"] == "Rectan
train_copy['count_floors_change'] = (train_copy['count_floors_post_eq'] - train_copy
train_copy['height_ft_change'] = (train_copy['height_ft_post_eq'] - train_copy['heig
```

```
In [28]: train_copy.drop(['count_floors_pre_eq', 'height_ft_pre_eq'], axis=1, inplace=True)
```

```
In [29]: remove_columns = ["legal_ownership_status", "land_surface_condition", "foundation_type
def dropColumns(res_train_copy, remove_columns):
    for i in remove_columns:
        res_train_copy.drop([i], axis = 1, inplace = True)
    return res_train_copy
```

```
In [30]: train_copy = dropColumns(train_copy, remove_columns)
```

```
In [31]: train_copy.shape
```

```
Out[31]: (631761, 48)
```

```
In [32]: train_copy['count_families'].fillna(train_copy['count_families'].mode()[0], inplace=True)
```

```
In [33]: train_one_hot = pd.get_dummies(train_copy)
train_one_hot.drop(["district_id_x", "vdcmun_id_x", "ward_id_x"], axis = 1, inplace = True)
```

Supervised Learning

Modeling, Training and Prediction

The dataset has been cleaned, pre-processed and analyzed for understanding the dataset. After such a process, and yet before coming to modeling, the dataset has to split up into two parts: Train and Test dataset. The training dataset is used to train the algorithm used to prepare an algorithm to comprehend. To learn and deliver results. It incorporates both input data and the desired output. The test data collection is utilized to assess how well the algorithm was prepared with the trained dataset.

By using the Supervised learning Algorithms to train the dataset and also to test, predictions were made as to the desired outcome. The system was able to split, train and test the dataset. Along with that, the feature importance was also given as the output where it had the percentage of these possibilities in occurring in the mere future datasets that would be added.

The aim for us is to predict the Damage grade assigned to the building after assessment of Earth Quake
The following are Supervised Algorithms which we used for prediction

```
In [34]: ▶ z_train = train_one_hot['damage_grade']  
train_one_hot.drop(['damage_grade'], axis = 1, inplace = True)
```

```
In [35]: ▶ from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(train_one_hot, z_train, test_size
```

Logistic Regresion

Logistic regression is named for the function used at the core of the method, the logistic function. The logistic function, also called the sigmoid function was developed by statisticians to describe properties of population growth in ecology, rising quickly and maxing out at the carrying capacity of the environment. It's an S-shaped curve that can take any real-valued number and map it into a value between 0 and 1, but never exactly at those limits. $1 / (1 + e^{-\text{value}})$ Where e is the base of the natural logarithms (Euler's number or the EXP() function in your spreadsheet) and value is the actual numerical value that you want to transform.

```
In [36]: > #importing required module
from sklearn.linear_model import LogisticRegression

#intializing model
logistic_reg = LogisticRegression(solver="liblinear")


#fitting the train data
logistic_reg.fit(X_train, y_train)

# predicting the data using test data
log_predicted = logistic_reg.predict(X_test)
probs = logistic_reg.predict_proba(X_test)

#finding the confusion matrix
conf_mat = confusion_matrix(y_true=y_test, y_pred=log_predicted)

# visualizing confusion matrix
sns.heatmap(conf_mat, annot=True, annot_kws={"size":16}, fmt="d", cbar=False, linewidths=1)
plt.title("Confusion matrix of the classifier", fontsize=14)
plt.ylabel("Actual label", fontsize=12)
plt.xlabel("log_predicted label", fontsize=12)
plt.show()

# printing the classifcation report
print("Understanding Classifciation:\n")
print(classification_report(y_test, log_predicted))
```



Understanding Classifciation:

	precision	recall	f1-score	support
1	1.00	0.70	0.82	12331
2	0.45	0.37	0.40	17165
3	0.40	0.42	0.41	24369
4	0.60	0.75	0.67	30591
5	1.00	0.95	0.97	41897
accuracy			0.70	126353
macro avg	0.69	0.64	0.66	126353
weighted avg	0.71	0.70	0.70	126353

Naive Bayes

Naive Bayes is a classification algorithm for binary (two-class) and multi-class classification problems. The technique is easiest to understand when described using binary or categorical input values. It is called naive Bayes or idiot Bayes because the calculation of the probabilities for each hypothesis are simplified to make their calculation tractable. Rather than attempting to calculate the values of each attribute value $P(d_1, d_2,$

d3|h), they are assumed to be conditionally independent given the target value and calculated as $P(d1|h) * P(d2|H)$ and so on. This is a very strong assumption that is most unlikely in real data, i.e. that the attributes do not interact. Nevertheless, the approach performs surprisingly well on data where this assumption does not hold.

Here our classification is for Multi-class

```
In [37]: > from sklearn.naive_bayes import GaussianNB
naive_bayes = GaussianNB()


#fitting the train data
naive_bayes.fit(X_train, y_train)

# predicting the data using test data
naive_bayes_predicted = naive_bayes.predict(X_test)
probs = naive_bayes.predict_proba(X_test)

#finding the confusion matrix
conf_mat = confusion_matrix(y_true=y_test, y_pred=naive_bayes_predicted)

# visualizing confusion matrix
sns.heatmap(conf_mat, annot=True, annot_kws={"size":16}, fmt="d", cbar=False, linewidths=1)
plt.title("Confusion matrix of the classifier", fontsize=14)
plt.ylabel("Actual label", fontsize=12)
plt.xlabel("naive_bayes_predicted label", fontsize=12)
plt.show()

# printing the classification report
print("Understanding Classification:\n")
print(classification_report(y_test, naive_bayes_predicted))
```



Understanding Classification:

	precision	recall	f1-score	support
1	0.69	0.83	0.75	12331
2	0.44	0.17	0.25	17165
3	0.37	0.33	0.35	24369
4	0.54	0.76	0.63	30591
5	0.99	0.95	0.97	41897
accuracy			0.67	126353
macro avg	0.61	0.61	0.59	126353
weighted avg	0.66	0.67	0.65	126353

Decision Tree

Decision trees classify instances by sorting them down the tree from the root to some leaf node, which provides the classification of the instance. An instance is classified by starting at the root node of the tree, testing the attribute specified by this node, then moving down the tree branch corresponding to the value of the attribute as shown in the above figure. This process is then repeated for the subtree rooted at the new node

```
In [38]: ▶ #importing required module
from sklearn.tree import DecisionTreeClassifier
decision_tree = DecisionTreeClassifier(random_state=0, class_weight="balanced")


#fitting the train data
decision_tree.fit(X_train, y_train)

# predicting the data using test data
decision_tree_predicted = decision_tree.predict(X_test)
probs = decision_tree.predict_proba(X_test)

#finding the confusion matrix
conf_mat = confusion_matrix(y_true=y_test, y_pred=decision_tree_predicted)

# visualizing confusion matrix
sns.heatmap(conf_mat, annot=True, annot_kws={"size":16}, fmt="d", cbar=False, linewidths=1)
plt.title("Confusion matrix of the classifier", fontsize=14)
plt.ylabel("Actual label", fontsize=12)
plt.xlabel("decision_tree_predicted label", fontsize=12)
plt.show()

# printing the classification report
print("Understanding Classification:\n")
print(classification_report(y_test, decision_tree_predicted))
```



Understanding Classification:

	precision	recall	f1-score	support
1	0.76	0.78	0.77	12331
2	0.37	0.39	0.38	17165
3	0.38	0.38	0.38	24369
4	0.58	0.55	0.56	30591
5	0.96	0.96	0.96	41897
accuracy			0.65	126353
macro avg	0.61	0.61	0.61	126353
weighted avg	0.65	0.65	0.65	126353

RandomForest

Random Forest is a popular machine learning algorithm that belongs to the supervised learning technique. It can be used for both Classification and Regression problems in ML. It is based on the concept of ensemble learning, which is a process of combining multiple classifiers to solve a complex problem and to improve the performance of the model.

As the name suggests, "Random Forest is a classifier that contains a number of decision trees on various subsets of the given dataset and takes the average to improve the predictive accuracy of that dataset." Instead of relying on one decision tree, the random forest takes the prediction from each tree and based on the majority votes of predictions, and it predicts the final output.

```
In [39]: > #importing required module
from sklearn.ensemble import RandomForestClassifier
random_forest = RandomForestClassifier(class_weight="balanced_subsample", random_state=42)


#fitting the train data
random_forest.fit(X_train, y_train)

# predicting the data using test data
random_forest_predicted = random_forest.predict(X_test)
probs = random_forest.predict_proba(X_test)

#finding the confusion matrix
conf_mat = confusion_matrix(y_true=y_test, y_pred=random_forest_predicted)

# visualizing confusion matrix
sns.heatmap(conf_mat, annot=True, annot_kws={"size":16}, fmt="d", cbar=False, linewidths=1)
plt.title("Confusion matrix of the classifier", fontsize=14)
plt.ylabel("Actual label", fontsize=12)
plt.xlabel("random_forest_predicted label", fontsize=12)
plt.show()

# printing the classification report
print("Understanding Classification:\n")
print(classification_report(y_test, random_forest_predicted))
```



Understanding Classification:

	precision	recall	f1-score	support
1	0.82	0.78	0.80	12331
2	0.43	0.43	0.43	17165
3	0.41	0.40	0.41	24369
4	0.61	0.66	0.64	30591
5	0.99	0.95	0.97	41897
accuracy			0.69	126353
macro avg	0.65	0.65	0.65	126353
weighted avg	0.69	0.69	0.69	126353

K-Nearest Neighbours

K-Nearest Neighbors is a machine learning technique and algorithm that can be used for both regression and classification tasks. K-Nearest Neighbors examines the labels of a chosen number of data points surrounding a target data point, in order to make a prediction about the class that the data point falls into. K-Nearest Neighbors (KNN) is a conceptually simple yet very powerful algorithm, and for those reasons, it's

one of the most popular machine learning algorithms. Let's take a deep dive into the KNN algorithm and see exactly how it works. Having a good understanding of how KNN operates will let you appreciate the best and worst use cases for KNN.

```
In [40]: ▶ #imported required model
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()

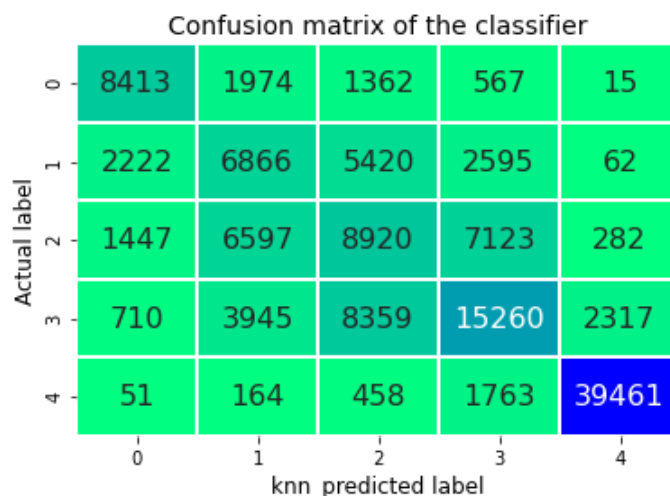
#fitting the train data
knn.fit(X_train, y_train)

# predicting the data using test data
knn_predicted = knn.predict(X_test)
probs = knn.predict_proba(X_test)

#finding the confusion matrix
conf_mat = confusion_matrix(y_true=y_test, y_pred=knn_predicted)

# visualizing confusion matrix
sns.heatmap(conf_mat, annot=True, annot_kws={"size":16}, fmt="d", cbar=False, linewidths=1)
plt.title("Confusion matrix of the classifier", fontsize=14)
plt.ylabel("Actual label", fontsize=12)
plt.xlabel("knn_predicted label", fontsize=12)
plt.show()

# printing the classification report
print("Understanding Classification:\n")
print(classification_report(y_test, knn_predicted))
```



Understanding Classification:

	precision	recall	f1-score	support
1	0.66	0.68	0.67	12331
2	0.35	0.40	0.37	17165
3	0.36	0.37	0.36	24369
4	0.56	0.50	0.53	30591
5	0.94	0.94	0.94	41897
accuracy			0.62	126353
macro avg	0.57	0.58	0.57	126353
weighted avg	0.63	0.62	0.63	126353

Confusion Matrix

A confusion matrix is a technique for summarizing the performance of a classification algorithm.

Classification accuracy alone can be misleading if you have an unequal number of observations in each class or if you have more than two classes in your dataset.

Calculating a confusion matrix can give you a better idea of what your classification model is getting right and what types of errors it is making.

Confusion matrix is square matrix of size n where n is the number of values for targeted variable

Let us consider a confusion matrix $A_{n \times n}$

$A[i,j]$ indicates the number of times i is predicted as j

In our case $n=5$ 0th row - grade 1, 1st row - grade 2, 2nd row - grade 3, 3rd row - grade 4, 4th row - grade 5

Classification Report

True Positives (TP) -

These are the correctly predicted positive values which means that the value of actual class is yes and the value of predicted class is also yes.

True Negatives (TN) -

These are the correctly predicted negative values which means that the value of actual class is no and value of predicted class is also no.

False Positives (FP) –

When actual class is no and predicted class is yes.

False Negatives (FN) –

When actual class is yes but predicted class is no.

Accuracy -

Accuracy is the most intuitive performance measure and it is simply a ratio of correctly predicted observation to the total observations. One may think that, if we have high accuracy then our model is best. Yes, accuracy is a great measure but only when you have symmetric datasets where values of false positive and false negatives are almost same. Therefore, you have to look at other parameters to evaluate the performance of your model. For our model, we have got 0.803 which means our model is approx. 80% accurate.

$$\text{Accuracy} = \frac{TP+TN}{TP+FP+FN+TN}$$

Precision -

Precision is the ratio of correctly predicted positive observations to the total predicted positive observations. The question that this metric answers is of all passengers that labeled as survived, how many actually survived? High precision relates to the low false positive rate. We have got 0.788 precision which is pretty good.

$$\text{Precision} = \frac{TP}{TP+FP}$$

Recall (Sensitivity) -

Recall is the ratio of correctly predicted positive observations to the all observations in actual class - yes. The question recall answers is: Of all the passengers that truly survived, how many did we label? We have got recall of 0.631 which is good for this model as it's above 0.5.

Recall = TP/TP+FN

F1 score -

F1 Score is the weighted average of Precision and Recall. Therefore, this score takes both false positives and false negatives into account. Intuitively it is not as easy to understand as accuracy, but F1 is usually more useful than accuracy, especially if you have an uneven class distribution. Accuracy works best if false positives and false negatives have similar cost. If the cost of false positives and false negatives are very different, it's better to look at both Precision and Recall. In our case, F1 score is 0.701.

F1 Score = $2 * (\text{Recall} * \text{Precision}) / (\text{Recall} + \text{Precision})$

Accuracy Calculation

```
In [41]: ▶ score_lr=accuracy_score(y_test, log_predicted)
          score_nb=accuracy_score(y_test, naive_bayes_predicted)
          score_dt=accuracy_score(y_test, decision_tree_predicted)
          score_rf=accuracy_score(y_test, random_forest_predicted)
          score_knn=accuracy_score(y_test, knn_predicted)

In [42]: ▶ scores = [score_lr,score_dt,score_nb,score_knn,score_rf]
          percentscore=[i*100 for i in scores]

In [43]: ▶ algorithms = ["Logistic Regression","Decision Tree","Naive Bayes","K-Nearest Neighbors"]
          for i in range(len(algorithms)):
              print("The accuracy score achieved using "+algorithms[i]+" is: "+str(percentscore[i]))
```

```
The accuracy score achieved using Logistic Regression is: 69.55671808346457%
The accuracy score achieved using Decision Tree is: 65.2283681432178%
The accuracy score achieved using Naive Bayes is: 66.78195214992917%
The accuracy score achieved using K-Nearest Neighbors is: 62.45993367787073%
The accuracy score achieved using Random Forest is: 68.86183944979541%
```

Accuracy Table

```
In [44]: ▶ from tabulate import tabulate
          data=[]
          for i in range(len(algorithms)):
              list=[algorithms[i],percentscore[i]]
              data.append(list)
          print(tabulate(data, headers=["Algorithm","Accuracy"]))
```

Algorithm	Accuracy
Logistic Regression	69.5567
Decision Tree	65.2284
Naive Bayes	66.782
K-Nearest Neighbors	62.4599
Random Forest	68.8618

Accuracy Graph

```
In [45]: ▶ sns.set(rc={'figure.figsize':(9,6)})  
plt.xlabel("Algorithms")  
plt.ylabel("Accuracy score")  
  
sns.barplot(algorithms,percentscore)
```

Out[45]: <AxesSubplot:xlabel='Algorithms', ylabel='Accuracy score'>

